


# Local Search Algorithms

# Local Search Algorithms

- Local search algorithms operate by searching from a start state to neighbouring states, without keeping track of the paths, nor the set of states that have been reached.
- That means they are not systematic—they might never explore a portion of the search space where a solution actually resides.
- However, they have two key advantages:
  - (1) they use very little memory; and
  - (2) they can often find reasonable solutions in large or infinite state spaces for which systematic algorithms are unsuitable.
- Local search algorithms can also solve **optimization problems**, in which the aim is to find the best state according to an **objective function**.



**Figure 4.1** A one-dimensional state-space landscape in which elevation corresponds to the objective function. The aim is to find the global maximum.

# Hill-climbing search

- It keeps track of one current state and on each iteration moves to the neighbouring state with highest value
  - that is, it heads in the direction that provides the steepest ascent.
- It terminates when it reaches a “peak” where no neighbour has a higher value.
- Hill climbing does not look ahead beyond the immediate neighbours of the current state.

- Each point (state) in the landscape has an “elevation,” defined by the value of the objective function.
- If elevation corresponds to an objective function then the aim is to find the highest peak—a ***global maximum***—and we call the process **hill climbing**.
- If elevation corresponds to cost, then the aim is to find the lowest valley—a ***global minimum***—and we call it **gradient descent**.

```
function HILL-CLIMBING(problem) returns a state that is a local maximum  
    current  $\leftarrow$  problem.INITIAL  
    while true do  
        neighbor  $\leftarrow$  a highest-valued successor state of current  
        if VALUE(neighbor)  $\leq$  VALUE(current) then return current  
        current  $\leftarrow$  neighbor
```

**Figure 4.2** The hill-climbing search algorithm, which is the most basic local search technique. At each step the current node is replaced by the best neighbor.

# Hill Descent: Example (Minimizing h)

---

start

3	1	2
4	5	8
6	7	

$h_{oop} = 5$

goal

	1	2
3	4	5
6	7	8

# Hill Descent: Example (Minimizing h)

---

start


3	1	2
4	5	8
6	7	


$h_{oop} = 5$

goal

	1	2
3	4	5
6	7	8

# Hill Descent: Example (Minimizing $h$ )






(a)

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	14	13	16	13	16
14	14	17	15	14	16	16	16
17	14	16	18	15	15	14	15
18	14	14	15	15	14	14	16
14	14	13	17	12	14	12	18


(b)

**Figure 4.3** (a) The 8-queens problem: place 8 queens on a chess board so that no queen attacks another. (A queen attacks any piece in the same row, column, or diagonal.) This position is almost a solution, except for the two queens in the fourth and seventh columns that attack each other along the diagonal. (b) An 8-queens state with heuristic cost estimate  $h=17$ . The board shows the value of  $h$  for each possible successor obtained by moving a queen within its column. There are 8 moves that are tied for best, with  $h=12$ . The hill-climbing algorithm will pick one of these.

# Reduce the number of conflicts



$$h = 5$$



- Hill climbing is sometimes called greedy local search because it grabs a good neighbor state without thinking ahead about where to go next.
- Hill climbing can get stuck for any of the following reasons:
  - Local maxima: A local maximum is a peak that is higher than each of its neighboring states but lower than the global maximum.
  - Ridges: Ridges result in a sequence of local maxima that is very difficult for greedy algorithms to navigate.
  - Plateaus: A plateau is a flat area of the state-space landscape

# Types of Hill Climbing Algorithms

- **Stochastic hill climbing:** chooses at random from among the uphill moves; the probability of selection can vary with the steepness of the uphill move.
- **First-choice hill climbing:** implements stochastic hill climbing by generating successors randomly until one is generated that is better than the current state.
- **Random-restart hill climbing:** It conducts a series of hill-climbing searches from randomly generated initial states, until a goal is found.

- A hill-climbing algorithm that never makes “downhill” moves toward states with lower value (or higher cost) is always vulnerable to getting stuck in a local maximum.
- In contrast, a purely random walk that moves to a successor state without concern for the value will eventually stumble upon the global maximum, but will be extremely inefficient.
- Combination of hill climbing with a random walk in a way that yields both efficiency and completeness.

# Simulated annealing

- In metallurgy, annealing is the process used to temper or harden metals and glass by heating them to a high temperature and then gradually cooling them, thus allowing the material to reach a low-energy crystalline state.
- Simulated-annealing solution is to start by shaking hard (i.e., at a high temperature) and then gradually reduce the intensity of the shaking (i.e., lower the temperature).
- Instead of picking the best move, however, it picks a random move. If the move improves the situation, it is always accepted.
- Otherwise, the algorithm accepts the move with some probability less than 1.

- Basic idea:
  - Allow “bad” moves occasionally, depending on “temperature”
  - High temperature => more bad moves allowed, shake the system out of its local minimum
  - Gradually reduce temperature according to some schedule

# Simulated annealing algorithm

```
function SIMULATED-ANNEALING(problem,schedule) returns a state
    current ← problem.initial-state
    for t = 1 to ∞ do
        T ← schedule(t)
        if T = 0 then return current
        next ← a randomly selected successor of current
        ΔE ← next.value – current.value
        if ΔE > 0 then current ← next
        else current ← next only with probability  $e^{\Delta E/T}$ 
```

# Simulated Annealing

- The probability decreases exponentially with the “badness” of the move—the amount  $\Delta E$  by which the evaluation is worsened.
- The probability also decreases as the “temperature”  $T$  goes down: “bad” moves are more likely to be allowed at the start when  $T$  is high, and they become more unlikely as  $T$  decreases.
- If the schedule lowers  $T$  to 0 slowly enough, then a property of the Boltzmann distribution,  
$$e^{\Delta E / T}$$
- Is that all the probability is concentrated on the global maxima, which the algorithm will find with probability approaching 1.

# Simulated Annealing


- Is this convergence an interesting guarantee?
- Sounds like magic, but reality is:
  - The more downhill steps you need to escape a local optimum, the less likely you are to ever make them all in a row
  - “Slowly enough” may mean exponentially slowly
  - Random restart hill climbing also converges to optimal state...
- Simulated annealing and its relatives are a key workhorse in VLSI layout and other optimal configuration problems

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  current  $\leftarrow$  problem.INITIAL
  for t = 1 to  $\infty$  do
    T  $\leftarrow$  schedule(t)
    if T = 0 then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}(\textit{current}) - \text{VALUE}(\textit{next})$ 
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{-\Delta E/T}$ 
```


**Figure 4.5** The simulated annealing algorithm, a version of stochastic hill climbing where some downhill moves are allowed. The *schedule* input determines the value of the “temperature” *T* as a function of time.

# Local beam search

- Basic idea:
  - $K$  copies of a local search algorithm, initialized randomly
  - For each iteration
    - Generate ALL successors from  $K$  current states
    - Choose best  $K$  of these to be the new current states
- Why is this different from  $K$  local searches in parallel?
  - The searches **communicate**! “Come over here, the grass is greener!”
- What other well-known algorithm does this remind you of?
  - Evolution!




# Genetic algorithms



- Genetic algorithms use a natural selection metaphor
  - Resample  $K$  individuals at each step (selection) weighted by fitness function
  - Combine by pairwise crossover operators, plus mutation to give variety

# Example: N-Queens



- Does crossover make sense here?
- What would mutation be?
- What would a good fitness function be?

# Summary

---

- Many configuration and optimization problems can be formulated as local search
- General families of algorithms:
  - Hill-climbing, continuous optimization
  - Simulated annealing (and other stochastic methods)
  - Local beam search: multiple interaction searches
  - Genetic algorithms: break and recombine states

Many machine learning algorithms are local searches